
Azure

# Azure SQL Database

Grow Data Skills

Azure SQL Database is a fully managed relational database service by Microsoft. It is based on the SQL Server database engine and offers high availability, scalability, and security for cloud applications.

**Key Features:**

- ***Fully Managed:*** Automated patching, backups, and monitoring.
- ***Scalable***: Handles dynamic workloads with elastic pools and autoscaling options.
- ***High Availability***: Built-in replication and geo-redundancy.
- ***AI-Driven Performance Optimization***: Intelligent query tuning and indexing.
- ***Security & Compliance***: Advanced data protection with built-in features like Transparent Data Encryption (TDE), Always Encrypted, and Dynamic Data Masking.

**Use Cases:**

- Web applications
- Business applications
- Analytics and reporting

# Azure SQL Database Deployment Models

Grow Data Skills

Azure SQL Database offers several deployment models to fit different needs:

**Single Database:**

- ***Overview***: Best suited for applications that require a dedicated and isolated database.
- ***Use Case:*** Independent application databases with predictable workloads.

**Elastic Pool:**

- ***Overview***: Multiple databases with varying and unpredictable usage patterns share resources within a pool.
- ***Use Case***: SaaS applications hosting databases for multiple tenants.

**Managed Instance:**

- ***Overview***: Provides a fully managed instance with near 100% compatibility with SQL Server. Ideal for migrating on-premises SQL Server workloads to Azure.
- ***Use Case***: Applications requiring SQL Server features like cross-database queries, linked servers, and SQL Agent.

**Comparison:**

- ***Single Database***: Ideal for standalone apps.
- ***Elastic Pool***: Cost-effective for managing multiple databases.
- ***Managed Instance***: Great for complex migrations needing high SQL Server compatibility.

# Pricing Model - Understanding DTU and vCore

Grow Data Skills

Azure SQL Database offers two main pricing models based on different performance measures:

**DTU-Based Model (Database Transaction Units):**

- ***Overview***: Combines CPU, memory, and read/write operations into a single performance measure (DTU).
- ***Use Case***: Simple to use for small to medium workloads.
- ***Service Tiers:***
  - Basic: Suitable for small databases with minimal traffic.
  - Standard: For general-purpose workloads with moderate traffic.
  - Premium: For high-performance and mission-critical applications.

**vCore-Based Model (Virtual Cores):**

- ***Overview***: Provides more control over individual resources like CPU, memory, and storage.
- ***Use Case***: Better for enterprise workloads with predictable performance needs.
- ***Service Tiers:***
  - General Purpose: Balanced option for most business applications.
  - Business Critical: For high-transaction workloads requiring low latency.
  - Hyperscale: For very large databases with auto-scaling capabilities.

**Comparison:**

- ***DTU Model:*** Simpler, bundled pricing based on performance levels.
- ***vCore Model***: More granular control, customizable options for CPU, memory, and storage.

***Choosing Between DTU and vCore:***

- ***DTU***: Best for simpler scenarios with less fine-tuning required.
- ***vCore***: Offers flexibility and is better suited for high-complexity workloads.

# Azure Key-Vault

Grow Data Skills

Azure Key Vault is a cloud service that provides secure storage and management of sensitive information such as API keys, passwords, certificates, and cryptographic keys. It helps in safeguarding access to secrets and ensures compliance with security standards.

**Key Features:**

- **Secrets Management:** Store and manage sensitive data like passwords, connection strings, and API keys securely.
- **Key Management:** Manage cryptographic keys used for encryption/decryption. It supports both software-protected and hardware security module (HSM) keys.
- **Certificate Management:** Securely store and manage digital certificates, automate certificate renewal, and reduce the risk of certificate expiry.
- **Access Control:** Provides granular access management using Azure Active Directory (AAD). You can control who can access and manage your secrets and keys.

**Use Cases:**

- **Secure Storage:** Centralized management for application secrets and credentials.
- **Encryption**: Manage keys for encrypting data in Azure Storage, SQL Database, and other Azure services.
- **Automated Certificate Management**: Automate certificate deployment and renewal for web apps and services.

**Benefits:**

- **Enhanced Security:** Secrets and keys are stored securely with role-based access control (RBAC).
- **Compliance**: Helps meet regulatory requirements by managing encryption keys in a compliant manner.
- **Simplified Key Rotation**: Automate key renewal and reduce the risk associated with manual management.

# Azure Data Lake Storage

Grow Data Skills

Azure Data Lake Storage (ADLS) is a highly scalable and secure data lake solution for big data analytics. It combines the flexibility and scalability of a data lake with the enterprise-grade security features of Azure, allowing you to store and analyze large volumes of structured, semi-structured, and unstructured data.

**Key Features:**

- **Scalability:** Store petabytes of data with high throughput for large-scale analytics.
- **Hierarchical Namespace:** Organize and manage data in directories and subdirectories, enabling file-level operations like rename and move.
- **Built-In Security:** Provides enterprise-grade security with access control lists (ACLs), role-based access control (RBAC), and integration with Azure Active Directory (AAD).
- **Cost-Effective Storage Tiers:** Multiple storage tiers (hot, cool, and archive) to optimize costs based on data access patterns.

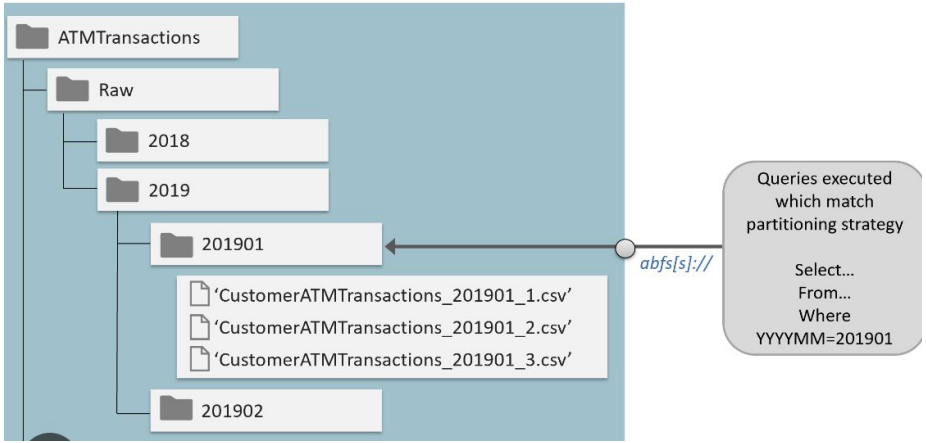

# Azure Event Hub

Grow Data Skills

Azure Event Hub is a fully managed, real-time data ingestion service that can stream millions of events per second. It is designed for high-throughput data streaming scenarios, such as telemetry data from IoT devices, application logs, or large-scale event processing systems. Event Hub serves as the "front door" for event ingestion, allowing downstream systems to process and analyze data in real-time.

**Architecture Components:**

- **Event Producers:** These are the sources generating events, such as applications, IoT devices, or services. Producers send events to Event Hub using the AMQP, HTTP, or Kafka protocols.
- **Event Hub Namespace:** A container for managing multiple Event Hubs. It acts as a logical grouping for all Event Hubs within an organization, simplifying management and resource organization.
- **Event Hub (Event Streams):** This is the main entity that receives and stores events. It can be seen as a highly scalable event stream divided into multiple partitions.
- **Partitions:** Event Hubs are divided into partitions, which are parallel streams within the hub. They help scale the event processing by enabling multiple consumers to read data in parallel. Each partition is ordered and allows for concurrent processing.
- **Consumer Groups:** A view (or cursor) of the event stream that is isolated from other views. Consumer groups allow multiple independent readers (e.g., analytics applications, monitoring tools) to process the events from the same partition independently without interfering with each other.
- **Event Consumers:** The receivers or processors that read the data from Event Hub. Consumers can be real-time processing systems (like Azure Stream Analytics, Apache Spark, or custom applications) that analyze, transform, and store the data.

# Azure Event Hub

Grow Data Skills

- **Event Hub Capture:** This feature allows automatic capturing of event data into Azure Blob Storage or Azure Data Lake Storage for batch processing or archiving. It's useful for data retention and batch analysis scenarios.

- **Throughput Units (TU) / Processing Units (PU):**
  - ***Throughput Units (for Standard Tier)***: Control the capacity of the Event Hub. A single TU allows for up to 1 MB per second of ingress (incoming data) and 2 MB per second of egress (outgoing data).
  - ***Processing Units (for Dedicated Tier)***: Dedicated clusters that allow for more customization and scaling needs for enterprise scenarios.
- **Offsets:** Act as pointers or bookmarks for event consumers, allowing them to start reading data from a specific point in the event stream, ensuring they pick up where they left off.
- **Checkpointing:** Checkpoints are used by consumers to keep track of their progress in reading data from partitions. It ensures that if a consumer restarts, it doesn't reprocess the same events.
-

**Event Hub Tiers:**

- ***Basic Tier:*** Suitable for smaller workloads with fewer requirements for features like consumer groups and throughput scaling.
- ***Standard Tier***: Offers more features like multiple consumer groups, throughput units, and Event Hub Capture.
- ***Dedicated Tier***: Provides a dedicated environment with processing units for large-scale, enterprise-grade streaming needs.

**Key Use Cases:**

- Real-time telemetry processing (e.g., IoT sensors)
- Log and event data streaming
- Analytics pipelines
- Event-driven microservices

# Azure Event Hub



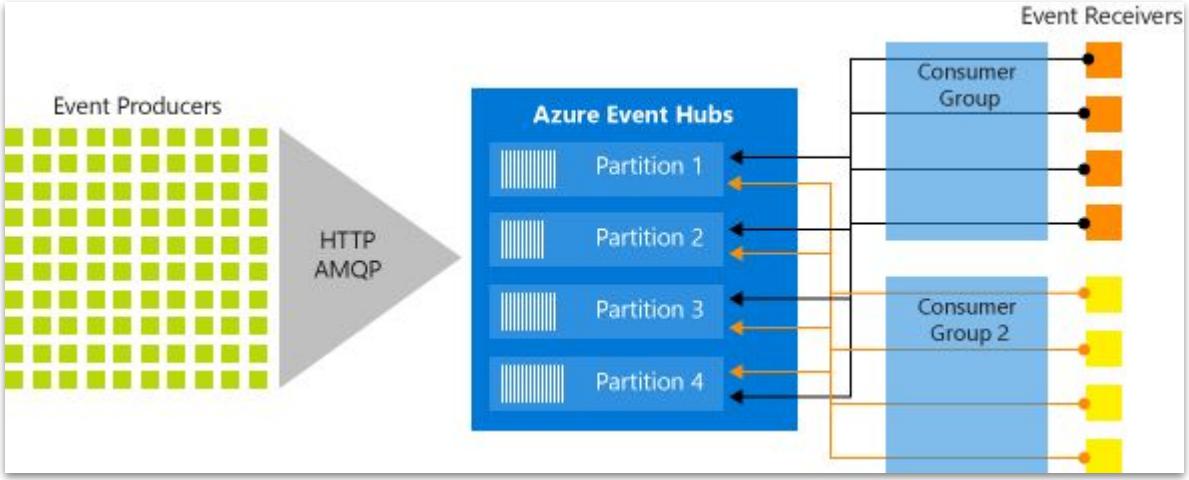


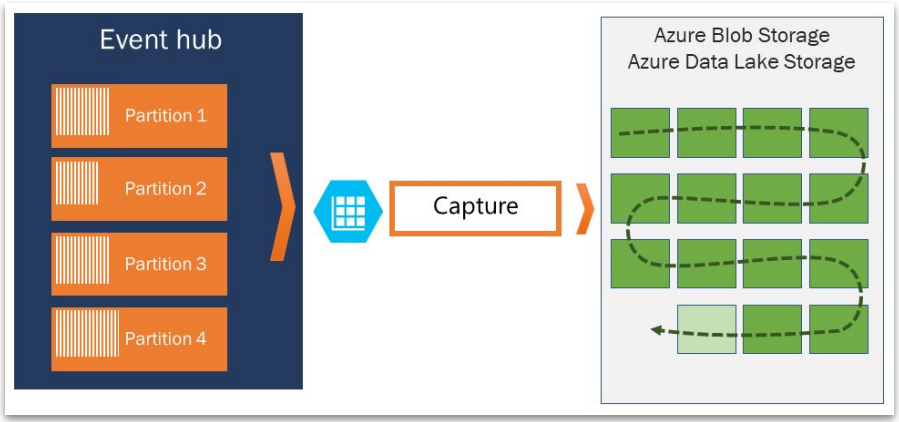

# Difference between Azure Event Hub and Kafka

Grow Data Skills

1. **Deployment & Management:**
   - **Event Hub:** Fully managed PaaS offering on Azure. Microsoft handles infrastructure, scaling, patching, and maintenance. Ideal for cloud-native applications with minimal operational overhead.
   - **Kafka:** Requires self-managed deployment (on-premises or in the cloud). Users handle cluster setup, scaling, broker management, and upgrades. Apache Kafka also has managed offerings like Confluent Cloud, but with additional costs.
2. **Data Ingestion & Throughput:**
   - **Event Hub:** Uses throughput units (TU) or processing units (PU) for capacity planning. Each TU provides 1 MB/s ingress and 2 MB/s egress. Automatic scaling within the predefined capacity; simpler but less flexible.
   - **Kafka:** Allows fine-grained control over partitions, replication, and broker configurations, enabling custom scaling strategies. Kafka's performance tuning offers more flexibility for high-throughput scenarios.
3. **Partitioning & Message Ordering:**
   - **Event Hub:** Partitions messages based on a partition key for message ordering within partitions. The number of partitions is defined at creation and cannot be changed later, making reconfiguration complex.
   - **Kafka:** Provides more flexibility in partitioning strategies. Partitions can be scaled dynamically, allowing more control over message distribution and ordering. Kafka supports leader election and rebalancing across partitions.
4. **Message Retention & Storage:**
   - **Event Hub:** Retention is time-based (up to 90 days for standard tier) and supports limited message replay. Designed more for short-term event streaming and real-time analytics use cases.
   - **Kafka:** Supports both time-based and size-based retention policies, allowing long-term data retention. Kafka's log-based storage is more suited for scenarios where replay and long-term data archiving are critical.
5. **Integration & Ecosystem:**
   - **Event Hub:** Natively integrates with Azure services (e.g., Stream Analytics, Azure Functions, Data Lake). Event Hub also supports Kafka protocol, allowing Kafka clients to connect directly. Best suited for enterprises already invested in the Azure ecosystem.
   - **Kafka:** Has a mature ecosystem with a wide range of connectors (via Kafka Connect) and stream processing tools like Kafka Streams, Flink, and Spark. Offers greater flexibility in integrating with diverse platforms and ecosystems.

# Azure Stream Analytics

Grow Data Skills

Azure Stream Analytics is a fully managed real-time analytics service designed to process large-scale streaming data from multiple sources, including Event Hubs, IoT Hubs, and Blob Storage.

**Key Features:**

- ***Real-Time Processing***: Analyze and act on data in motion with minimal latency.
- ***SQL-Like Query Language:*** Use a familiar SQL-like language for data transformation, filtering, and aggregation.
- ***Scalability and Flexibility***: Auto-scale to handle varying data volumes with built-in support for complex event processing.
- ***Integration with Azure Ecosystem:*** Seamlessly integrate with other Azure services like Power BI, Azure SQL Database, Event Hubs, and Data Lake.

**Common Use Cases:**

- ***IoT Monitoring***: Real-time insights from sensor data for predictive maintenance and anomaly detection.
- ***Clickstream Analysis***: Process user interaction data to drive personalization and recommendations.
- ***Fraud Detection***: Instant alerts and actions on suspicious transactions.

**How It Works:**

- ***Data Ingestion***: Input data from Event Hubs, IoT Hubs, or Blob Storage.
- ***Data Processing***: Define queries to filter, aggregate, and transform data in real-time.
- ***Output***: Send processed data to Azure SQL Database, Event Hubs, Power BI, Blob Storage, or custom endpoints.